

Belief-Based Utility and Signal Interpretation^{*}

Marta Kozakiewicz[†]

Job Market Paper

First version: October 17, 2020

This version: February 23, 2021

[Click here for the most recent version](#)

Abstract

People tend to overestimate their abilities and chances of success, even though inaccurate beliefs lead to costly mistakes. How can these beliefs persist in an environment with frequent feedback? I propose a new test of the hypothesis that people interpret favorable feedback to be more informative. Using experimental data, I provide the first causal evidence that the utility from beliefs affects one's perception of signal informativeness. To establish causality, I adopt a matching estimator approach and construct a counterfactual outcome of a subject who observes the same signal, but the signal is not affecting his belief-based utility. I find a strong and significant effect: subjects interpret favorable signals to be more informative due to changes in belief-based utility. The results cast a new light on the origins of overconfidence and illuminate mechanisms that perpetuate it in the face of feedback.

Keywords: overconfidence, belief formation, learning, experiment

JEL classification: C91, D83

^{*}The author gratefully acknowledges funding by the Deutsche Forschungsgemeinschaft (DFG, German Research Foundation) through CRC TR 224 (project A01).

[†]Bonn Graduate School of Economics; email: martkozakiewicz@gmail.com

1 Introduction

People tend to overestimate their abilities and chances of success, making costly mistakes as they hold on to their biased beliefs at the expense of accuracy. This tendency, commonly referred to as *overconfidence*, generates significant costs for both the individual and society¹. A long-standing question in behavioral economics is how it can persist in environments with frequent feedback. In this paper, I explore one possible explanation.² I consider an agent who does not know his ability and receives a signal that either reveals it or not. The agent forms beliefs about both his ability *and* informativeness of the signal. Importantly, he values his beliefs about his ability, so that any change in these beliefs directly affects his utility function (Brunnermeier and Parker, 2005; Caplin and Leahy, 2019; Kőszegi, 2006). I attempt to answer the following questions: Does the agent perceive a favorable signal to be more informative than an unfavorable one? Would he perceive the signal differently if the signal did not affect his utility function?

To this end, I designed a simple experiment in which participants learn about their performance in an IQ test.³ I randomly assigned subjects to one of two conditions. In a treatment condition, participants received a signal about their performance and reported their beliefs about the signal’s informativeness. In a control condition, subjects made a similar report for every possible signal realization. They faced *the same* decision but *without* receiving an actual signal. We observe payoff-maximizing reports absent signal-induced changes in beliefs and belief-based utility. Using observations from the control condition, I construct a counterfactual outcome for every subject in the treatment condition: what he would have reported if the signal had not affected his beliefs. The difference between reports in the treatment condition and the counterfactual reveals the extent of belief manipulation in response to favorable and unfavorable signals, and pins down a causal effect of signal valence on updating. Moreover, it informs us about the underlying mechanism by showing how belief-based utility affects signal interpretation.

¹Negative consequences of overconfidence include excessive selection into competitive environments (Camerer and Lovo, 1999; Niederle and Vesterlund, 2007), excessive trading (Barber and Odean, 2001), suboptimal investment decisions (Malmendier and Tate, 2005, 2008), and political polarization (Ortoleva and Snowberg, 2015).

²Other explanations that are similar to my work (as they consider motivated reasoning rather than cognitive processes) can be divided into three categories: information avoidance (see Golman et al., 2017, for a comprehensive literature review), selective recall (Chew et al., 2019; Huffman et al., 2019; Zimmermann, 2020), and asymmetric updating. The last point mentioned comes the closest to my work and I review it in detail in the following section.

³The experiment was pre-registered in the AEA RCT Registry (Registration Number AEARCTR-0006233). Details of the registration are provided in Appendix H.

The data support the hypothesis that changes in belief-based utility drive subjects’ perception of signal informativeness. I find a strong and significant effect, with a positive asymmetry: subjects perceive favorable signals to be more informative. In comparison to the counterfactual, participants reported a 10.6 percentage points higher probability of a favorable signal being entirely informative (a 27.9% increase). There is no difference after unfavorable signals. As a result, subjects in the treatment condition ended up with significantly *lower* payoffs compared to what they would have gotten if they were to decide without observing the signal realization. A striking conclusion of the experiment is that allowing subjects to acquire signals makes them worse-off (in monetary terms).

My study provides the first clear evidence of a causal effect of belief-based utility on signal interpretation. While the research on updating beliefs about ego-relevant traits has a long tradition (I review the literature in Section 2), establishing causality has always been challenging. One difficulty lies in introducing exogenous variation in “ego-relevance”: the way signals affect belief-based utility. Ideally, we would like subjects to receive the same feedback, but the feedback would have no *valence* – it would not be “positive” or “negative” in a sense that it would not bring participants additional belief-based utility. But how to separate feedback from its valence?

Previous work focused on comparing how people update their beliefs about some ego-relevant characteristic (e.g. one’s performance in an IQ test) and how they update beliefs about some ego-neutral parameter (e.g. performance of a robot).⁴ However, this comparison involves not only learning about ego-relevant and ego-neutral parameter, but also updating subjective beliefs, possibly multiple priors, and updating objective probabilities given by the experimenter. The experimental manipulation affects more than one aspect of the study undermining causal inference.

In this paper, I propose a novel experiment in which both the treatment and the control condition are based on the same subjective beliefs over the same ego-relevant characteristic. However, I introduce exogenous variation in how signals affect subjects’

⁴See, for instance, Coutts (2019), Eil and Rao (2011), and Möbius et al. (2014). One exception is a study by Buser et al. (2018), which compares how participants update beliefs about their performance in various tasks that differ in how relevant they are to the subject’s self-esteem. However, in their set-up, it is not possible to introduce exogenous variation in ego-relevance. Grossman and Owens (2012) propose a control condition in which participants learn about the test result of another subject. In this case, subjects update their subjective beliefs about an unknown, ego-neutral variable.

beliefs and their belief-based utility: in the control condition, a signal is not realized, hence it does not affect subjects’ beliefs nor their belief-based utility. Thereby, I separate feedback from its valence without changing other decision-relevant aspects of the design.⁵

The study was conducted in August 2020 in the BonnEconLab at the University of Bonn. In total, I collected data from 222 participants. The experiment consisted of several parts. Firstly, participants were given an IQ test and incentivized to do their best. After the test, they were asked to report their beliefs about their relative performance. Using an incentive compatible mechanism, I elicited subjective beliefs about one’s test score falling into the 1st, 2nd, ..., 10th decile of the score distribution. I referred to the deciles as “ranks”, with 1 denoting the highest and 10 denoting the lowest rank.

After the belief elicitation, we described the framework to the subjects as follows: “There are two boxes. Box 1 contains 10 balls with numbers 1 to 10 written on them (each number occurs exactly once). Box 2 contains 10 balls with the same number written on every one of them. That number is equal to your rank.” For example, if a subject’s rank is 4, Box 2 contains 10 balls with the number “4” written on them.

In the main task, one ball was randomly drawn from one of the boxes (either box could be selected with equal probability) and presented to the subject. After seeing the ball, the participant reported his beliefs about the event that the ball came from Box 2 (with his rank). The report was made by dividing 100 points between the two boxes. I incentivized truthful reporting with the Binarized Scoring Rule (Hossain and Okui, 2013). The method was explained to the participants and they were informed that their chances to win the highest reward were maximized when they divided their points in a way that corresponded to their true beliefs about the box. We explained in intuitive terms how one can arrive at a Bayesian update given one’s prior beliefs about the rank.

An ideal counterfactual to the treatment condition would include a subject who has the same prior belief distribution (or the same set of prior belief distributions if the agent had multiple priors) and observes the same signal, but the signal has no effect on his belief-based utility function. To come as close as possible to the ideal counterfactual, I designed a control condition, which I describe below.

⁵ Alternatively, one could introduce exogenous variation in signal precision. I decided to use a different approach to measure the full scope of belief distortion in response to signals and investigate the role of belief-based utility in asymmetric updating.

In the control condition, subjects do not see a ball being drawn, but are asked to report their beliefs about signal informativeness *ex ante*, for every possible signal realization. The procedure, known as the Strategy Method, is commonly used in experiments investigating strategic interactions in games (Brandts and Charness, 2009). To alleviate concerns about the non-comparability of the two treatments, I adopted procedures that specifically targeted the issues raised in the literature.⁶ I argue that a participant in the control condition faces the same decision as a subject in the treatment condition but without the signal affecting his beliefs and belief-based utility.

Note that, although the assignment of subjects to the treatment and control condition was random, the assignment of signals to subjects was not. Participants in the treatment condition make a report about one number that, with probability $\frac{1}{2}$, is their rank. In the control condition, every participant reports his beliefs about all 10 numbers. This leads to a correlation between the treatment status and numbers considered by subjects. To draw a causal inference, I construct a matching estimator. For every participant in the treatment condition, I construct a counterfactual outcome using all observations from the control regarding *the same number* as the one seen by the subject in the treatment condition. However, not all of these observations receive the same weight. Those participants in the control condition, whose true rank and prior belief distribution were closer to the rank and beliefs of the participant in the treatment condition, receive a higher weight.⁷ I interpret the counterfactual as what the subject would report if he were in the control condition.

⁶One concern raised in the experimental game theory literature is that players may gain a better understanding of the game if they are induced to think about the best strategies from the perspective of other players. One can imagine that considering every possible signal in the control condition could influence subjects' beliefs. I address this issue by presenting participants in the treatment condition with the screenshots from the control condition and asking them to think about every possible draw before they proceed to the main task. While only participants in the control condition are allowed to enter their choices, both groups are required to consider every signal realization. Moreover, I hope to alleviate another concern, the problem of framing the answers in the strategy method with the order of options, by randomizing the order of the signals presented to the subjects in both conditions.

⁷I follow Heckman et al. (1998) and estimate the weights using kernel regression. I use non-parametric regression, as linear regression may be inefficient in presence of nonlinearities. As a robustness check, I run a simple linear regression using all observations from the control condition (I report the results in Appendix C). The effect has the same order of magnitude, although it is significant only at 0.1 level.

Moreover, I consider two alternative matching specifications: a matching based solely on the prior belief distribution (in theory, it subsumes information that a subject has about his performance), and a matching based on one's rank and prior belief about the signal at hand. The results are very similar to those of the first specification (I report them in Appendix F).

The results based on matched data lend unequivocal support to the initial hypothesis. There is a significant difference in the reported probability of a signal being informative in the treatment condition compared to the counterfactual. The effect is entirely driven by a differential response to signals that are above and below one’s median prior belief. Favorable signals are believed to be 10.6 percentage points more likely to be informative (a 27.9% increase compared to the counterfactual). The effect strongly depends on the subject’s expectations. It is no longer present if a subject assigned zero prior probability to the state of the world indicated by the favorable signal. For favorable signals to which subjects assigned non-zero prior probability, the difference increases to 15.7 percentage points. In contrast, there is no difference in subjects’ reports after unfavorable signals.

Using subjects’ responses in questionnaires, I provide additional evidence to support my interpretation of the results as being driven by changes in belief-based utility. In the treatment condition, those participants who report experiencing hopelessness (a negative anticipatory emotion) tend to deviate more from the Bayesian benchmark. The effect is counteracted by the habitual use of emotion regulation strategies. Subjects who reported using more emotion regulation in their daily life tend to deviate less from Bayesian updating, even if they admit to feeling more hopeless. While only suggestive, the evidence supports the view that the treatment effect is stemming from the visceral, emotion-based reaction to signals that are indicative of a belief-based utility.

The paper is organized as follows. In the next section, I describe the relevant literature. In Section 3, I outline the experimental design. Section 4 presents the empirical results, and Section 5 describes the additional evidence. Section 6 concludes.

2 Literature Review

My work is based on the theoretical literature on overconfidence and belief formation. That literature postulates that people derive utility not only from physical outcomes but also from their beliefs about the current or future state (Brunnermeier and Parker, 2005; Caplin and Leahy, 2019; Kőszegi, 2006). The individual can choose his beliefs but faces a trade-off between their accuracy (necessary to take the optimal action) and

their desirability (a consequence of the non-monetary value beliefs bring to the agent).⁸ The tension is resolved by the agent manipulating his beliefs to the extent that he is not losing too much from action taken based on those beliefs.

The main difficulty in empirically investigating the belief-based component of the utility function is that not only we cannot observe *preferences* over different belief distributions but also, in opposition to the physical outcomes, we have limited information about the resulting *choices*, as we usually do not observe the choice set: all the distributions of beliefs the agent is choosing from. Given these difficulties, it is unsurprising that we rarely model belief formation as a choice made, more or less consciously, by the agent. In fact, most studies conceptualize belief formation as beliefs updating assuming that beliefs, well-defined and probabilistically quantified, follow a pre-specified set of rules, Bayesian updating being the prime example. The Bayesian approach, bolstered by axiomatic derivation justifying its position of a rational benchmark, became the most prevalent model of beliefs updating (Gilboa and Marinacci, 2016).

Although it is a good approximation to reality in some contexts, the Bayesian model seems to be less adequate in others. One such an example is a situation in which the decision-maker has a clear preference over the states of the world, as the in case of learning about an ego-relevant trait. Several studies demonstrated that agents significantly deviate from Bayes' rule when forming beliefs about their own intelligence or beauty (Buser et al., 2018; Coutts, 2019; Eil and Rao, 2011; Ertac, 2011; Grossman and Owens, 2012; Möbius et al., 2014; Schwardmann and Van der Weele, 2019). The main conclusion emerging from this strand of literature is that belief formation over ego-relevant characteristics significantly differs from learning about ego-neutral variables. At the same time, the direction of the effect and its magnitude vary across studies.

The design presented in this paper differs from these experiments in important ways. First of all, subjects in my study observed only one signal. I aimed at disentangling the

⁸Some behavioral studies emphasize the consumption value of beliefs (due to pleasant or unpleasant emotional reactions they tend to induce), others stress the importance of non-classical instrumental value including motivational value, signaling value, or value from serving as a commitment device (see Bénabou and Tirole, 2016 for a comprehensive review of the literature).

effect of attribution to noise from the way agents are aggregating information (which may also be affected by motivated reasoning, but is beyond the scope of the paper).⁹

Secondly, in my set-up, the signal is either perfectly informative or entirely uninformative, with equal probability known to the subjects. This allows us to control for the extent to which subjects “compress” probabilities towards 50%, the effect observed in updating about ego-neutral variables (Ambuehl and Li, 2018; Enke and Graeber, 2019).

Moreover, I use a richer state and signal space compared to the above-mentioned studies. To understand why it is important, imagine a participant who believes that he is in the 80th percentile of the IQ test score distribution. Receiving a coarser signal, e.g. a signal indicating that his score was above the median, would not influence his beliefs as it merely confirms what he already knows. However, if the signal was more precise, e.g. it revealed that his score was only in the 60th percentile, it would affect his beliefs and, according to my hypothesis, induce a stronger reaction.

The idea presented in this paper is related to research on emotions and decision-making (Lerner et al., 2015). One conclusion from the psychological literature is that emotions may influence decisions via changes in the content of thought, and vice versa. A similar hypothesis has been tested in a recent study of Engelmann et al. (2019) who investigate the impact of anxiety on wishful thinking. Using data from a carefully designed experiment, they show a causal effect of anticipatory anxiety on belief formation. Although I cannot argue about the causal impact of anticipatory emotions in my experiment, the suggestive evidence is in line with their findings.

3 Experimental Design

The experiment consisted of two parts and is outlined in Figure 1. In the first part, subjects completed an IQ test intended to assess their cognitive ability. The second part included the elicitation of prior and posterior beliefs and a stage in which subjects received signals (or considered every possible signal realization in the control condition).

I describe the procedures in detail in the following subsections.

⁹For that reason, my experiment is also related to the literature on self-serving attribution bias. It has been extensively studied by psychologists (see Mezulis et al., 2004, for a meta-analysis of the existing studies) and, more recently, by economists (Coutts et al., 2020; Hestermann and Yaouanq, 2020; Van den Steen, 2004). None of the studies, however, consider the counterfactual discussed in my paper.




Figure 1: The outline of the experiment.

3.1 IQ Test

In the first part of the experiment, I evaluated the subjects' cognitive ability using an IQ test.¹⁰ The test consisted of 29 standard logic questions and participants were asked to solve as many of them as possible in 10 minutes. Individual scores were calculated based on the number of correctly answered questions minus the number of incorrect answers, and subjects were paid 0.75 Euro for every point they obtained.

¹⁰I decided to use intelligence as a basis for the learning exercise for several reasons. Firstly, it is known that intelligence correlates strongly with educational achievement, success in the labor market, and income. Because of that, I expect people to care deeply about their cognitive ability. Therefore, IQ measure seems to be a good candidate for a genuine ego-relevant parameter. Secondly, the literature provides evidence that people have biased beliefs about their cognitive ability (with overconfidence prevailing among men), which suggests that learning about one's cognitive ability may be one of natural settings in which the mechanism is in play.

Participants were informed that their earnings from the IQ test will be added to their earnings from the remaining parts of the experiment and paid at the end of the session. They were also informed that, although they will receive the entire sum of money at the end of the study, they will not learn immediately the exact number of points they obtained in the IQ test, nor how much money they earned in each part. Participants were informed that their IQ test results and the details of their payoffs will be available to them in one week after the session. Every participant received a personal link to a website on which his individual information was posted one week later.¹¹

3.2 Belief Elicitation

At the beginning of the second part, participants were told that they have to complete 3 tasks, for which they can earn up to 12 Euro. They were informed that *one task* will be drawn at random at the end of the session, and they will be paid only for that task.

In the first task, I elicited subjects' beliefs about their test scores being in the 1st, 2nd, ..., 9th and 10th deciles of the distribution of the test scores of 300 participants who took the same test in the BonnEconLab in previous sessions. I introduced 10 "ranks", with Rank 1 denoting the highest rank (assigned to participants whose IQ test scores were higher than or equal to the test scores of 90 – 100% of all participants), and Rank 10 denoting the lowest rank (defined analogously). The first task was to allocate 100 points among the ranks in a way that reflects one's beliefs about the relative performance in the IQ test.

The screen-shot of the computer interface used by subjects is presented in Figure 2. Participants were allocating points by dragging blue arrows to selected positions. They were informed that they can move the arrows back and forth to correct their choices. The text below the scales informed a participant how many points are being allocated to a given rank and the allocation was immediately appearing on the graph to the right.

¹¹This procedure served two purposes. First of all, I wanted to minimize dynamic concerns (e.g. subjects may adopt overly pessimistic beliefs to prepare themselves for the arrival of "bad news"). Secondly, this feature of the design enables me to collect data on who decided to check the test results. I describe the data on information acquisition in Appendix I.

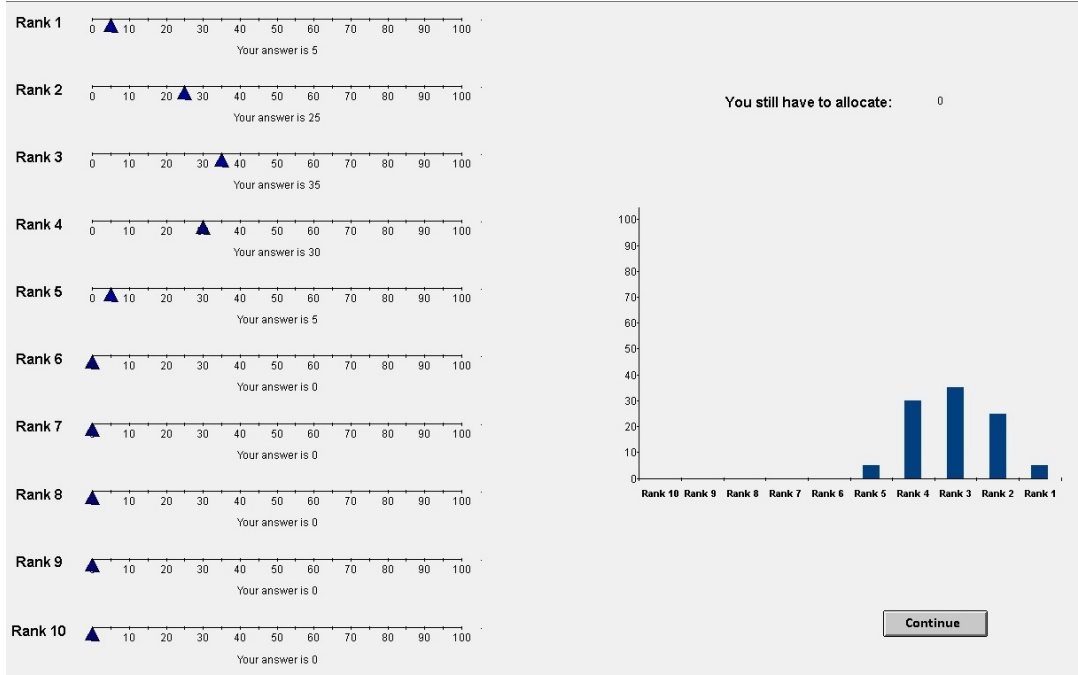


Figure 2: The screen-shot of the interface used by subjects in belief elicitation.

The number above the graph indicated how many points the participant still has to allocate before he can proceed to the next task.

To incentivize truthful reports, I used the Binarized Scoring Rule following Hos-sain and Okui (2013). The random variable X can take one of 10 values: $(1,0,\dots,0,0)$, $(0,1,\dots,0,0)$, ..., $(0,0,\dots,1,0)$, $(0,0,\dots,0,1)$; the position of 1 indicates in which decile subject's IQ test score fell. After receiving agent's report $x = (x_1, \dots, x_{10})$, where x_i denotes the share of points allocated to decile $i \in \{1, \dots, 10\}$, I observed his IQ test score in the k^{th} decile, and the agent won the prize if the QSR for multiple events,

$$s(x, k) = 2x_k - \sum_i x_i^2 + 1,$$

exceeded a uniformly drawn random variable with the support $[0, 2]$.

The formula was presented to the subjects in a simple way (avoiding mathematical notation). Importantly, I told participants the main implication of the method, that is,




Figure 3: The composition of the boxes of a person whose rank was 2.

the probability of getting a large prize (12 Euro) is maximized when they allocate their points in a way that reflects their beliefs about their rank.

I followed the same procedure during the second belief elicitation, after the signal stage (after the strategy method in the control condition). However, during the first belief elicitation, subjects were not aware that they will be asked to state their beliefs one more time.

3.3 The Signal Stage

After eliciting the prior beliefs, participants were given instructions for the second task. We explained the nature of the task in a simple language, using pictures and two illustrative examples. The task was framed in a neutral way and described as follows.

There are two boxes: Box 1 and Box 2. Each box contains 10 balls with numbers written on them. Box 1 contains balls with numbers from 1 to 10, and every number appears exactly once. The composition of the second box depends on the subject's rank in the IQ test. Box 2 contains 10 balls that all have one number written on them, and this number is equal to the individual rank. The composition of the boxes of a person assigned Rank 2 is presented in Figure 3.

For every participant, the computer program randomly selected one of the two boxes. Next, a ball was drawn from the selected box and displayed on the participant's screen.

The participant did not know which box the ball was drawn from, but he knew that either box can be selected with equal probability. After seeing the ball, he had to state his beliefs about the box selected by the computer.

I used the same incentive-compatible elicitation method as for the prior and posterior belief elicitations. Participants had 100 points to allocate between Box 1 and Box 2 in proportions that reflect their beliefs about the source of the signal, and were rewarded for the truthful report with a higher probability of getting a large prize (12 Euro).

Importantly, subjects were instructed how to arrive at the Bayesian posterior given one's prior belief distribution. I explained it with an example in two steps. Firstly, I demonstrated how a person should allocate her points after different signal realizations if she knew precisely her rank. Then, I showed how a person should allocate her points if she was not sure about her rank, but was assigning a certain probability to it.

Step 1: How should a person ranked 2 allocate her points if she knew for sure that her rank is 2, and saw a ball with a number "2" on it? There are 10-times as many balls with "2" in Box 2 as there are in Box 1, hence it is 10-times as likely that the ball came from the second box. Therefore, the person should allocate 9 points to Box 1, and 10-times as many, 90 points, to Box 2 (the remaining point should be allocated to the box with higher probability).

Step 2: What if a person did not know her true rank, but she believed that there is 30% chance that her rank is 2? The same logic applies to this case. One can visualize 30% chance as 3 out of 10 balls in Box 2 having a number "2" on them.¹² In this imaginary case, there are 3-times as many balls with the number "2" on them in Box 2 as in Box 1, implying an allocation of 25 points to Box 1 and 3-times as many (75 points) to Box 2.

The interface enabled subjects to split their points in desired proportions without calculating the respective ratios. The screen-shot of the interface used in the second task is presented in Figure 4. Crucially, the text below the scale informed subjects about their

¹²One reason why I decided to introduce 10 balls was the ease of exposition in a case when a person is uncertain about his rank.

current allocation and the ratio between points allocated to the two boxes. By moving the cursor, participants could choose the number of points corresponding to allocating x -times as many points to one of the boxes (with $x \in \{1, 1.1, \dots, 99\}$). The graph below was illustrating the current allocation.

Before proceeding to the signal stage, participants were required to answer a set of control questions, designed to check their understanding of the task (including the steps necessary for arriving at the Bayesian posterior). The control questions also pointed out the aspects that participants may have missed at the first reading, but were necessary to fully comprehend the task.

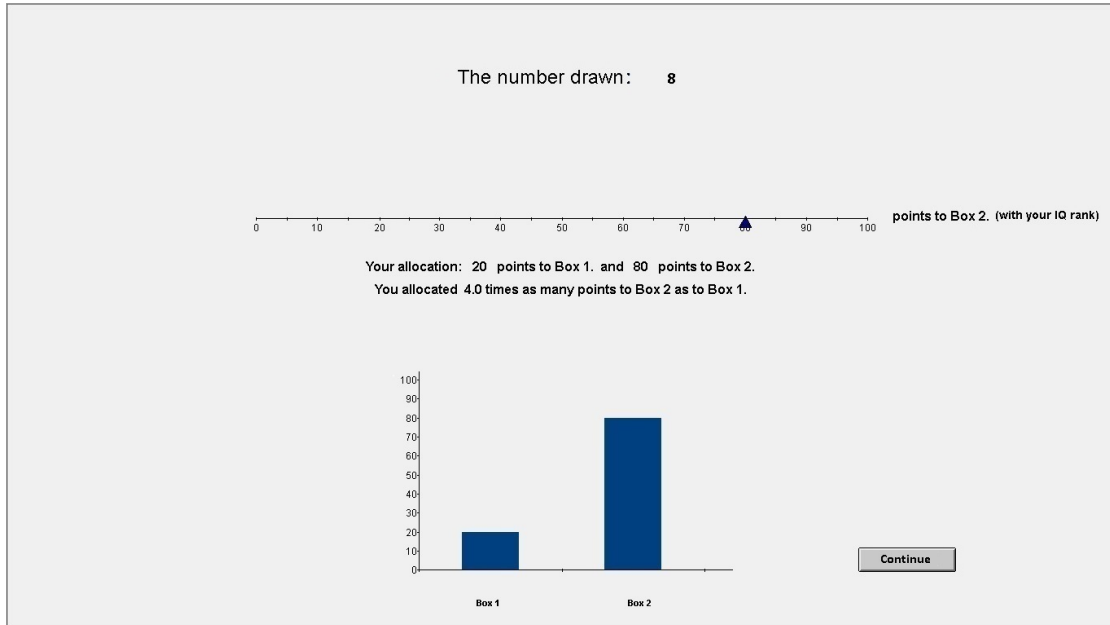


Figure 4: The screen-shot of the interface used in the second task (the signal stage).

3.4 Experimental Conditions

I introduced two experimental conditions: treatment and control. In the control condition, subjects did not see the number that was drawn but were asked to state their beliefs for every possible draw. The procedure, known as the Strategy Method, is commonly used in experiments investigating strategic interactions in games.

I informed participants in the control condition that the choices they are making are not entirely hypothetical. At the end of the session, one box was selected by the computer program and one ball was randomly drawn from the selected box. Subjects were paid as in the treatment condition, based on the decision that corresponded to the number drawn from the box. Note that the procedure is incentive-compatible as the probability of drawing any number is at least 5%.¹³

To alleviate concerns of the non-comparability of the two conditions, I adopted special procedures targeting the issues discussed in the literature. One concern raised in the experimental game theory literature is that players in the strategy method gain a better understanding of the game as a consequence of considering the problem from the point of view of different players. In my set-up, one can imagine that considering every possible signal realization may influence reported beliefs in the control condition.

For this reason, we asked the participants in the treatment condition to consider every possible signal realization *before* they saw the actual draw. Subjects were required to go through 10 slides, presented in random order, with the actual screen-shots of the interface displayed in the control condition. Participants were asked to contemplate a hypothetical decision in each slide before clicking on the button “Continue”, which appeared on the screen only after 15 seconds. While only subjects in the control condition were allowed to enter their choices, both groups were required to go through the task.

Another problem that may arise in the Strategy Method is framing the answers with the order of options. I addressed the issue by randomizing the order of the numbers displayed to a subject in the control condition, and the order of slides presented to participants in the treatment.

¹³However, if subjects were weighting the cost of cognitive effort against the expected payoff, they may exert less effort in the control condition. In this case, one would expect subjects to behave *less* rationally: their decisions would be characterized by a higher variance and they would end up further away from Bayesian update. This is the opposite of what I found.

3.5 Questionnaires

After each part of the experiment, I asked participants to fill in a 3-page questionnaire. The first set of questions, displayed on individual computer screens after the IQ test, included a short version of the Big-5 personality test (Gerlitz and Schupp, 2005) and the state-trait anxiety inventory STAI (Spielberger, 1983).

The Big-5 personality test was designed to measure personality along five dimensions: extroversion, conscientiousness, openness to experience, neuroticism, and agreeableness. The STAI measures the current state of anxiety and anxiety level as a personal characteristic. The second set of questions, answered by the participants after the main task, comprised the Emotion Regulation Questionnaire (Gross and John, 2003) and a subset of questions from the Achievement Emotions Questionnaire (Pekrun et al., 2011).

The Emotion Regulation Questionnaire was designed to assess the habitual use of two strategies commonly used to alter emotions. To alleviate the emotional impact of a situation, one may try to reinterpret it in a different way. This emotion regulation strategy, broadly referred to as *reappraisal*, relies on “applying mental models to the often ambiguous and incomplete information” (Uusberg et al., 2019). The second emotion regulation strategy, *suppression*, involves “inhibiting ongoing emotion-expressive behavior” (Gross and John, 1998, cited in Uusberg et al., 2019).

People differ in their use of reappraisal and suppression, and these differences have implications for their experiences of emotions, behavior in response to those emotions, and general well-being (Gross and John, 2003). The habitual use of the two strategies is measured by the degree to which subjects agree with particular statements, e.g. “I keep my emotions to myself” or “When I want to feel less negative emotion, I change the way I’m thinking about the situation”. I use the exact 10-item questionnaire developed by Gross and John (2003).

The Achievement Emotions Questionnaire was designed to measure *achievement emotions* (emotions that are directly linked to achievement activities or achievement outcomes) experienced by students in academic settings (Pekrun et al., 2011). I adopted

part of the questionnaire to measure the following test-related emotions: enjoyment, hope, pride, relief, anger, anxiety, shame, and hopelessness.

Participants in both conditions were asked to report what they felt *after* learning the nature of the task, but *before* they saw the number(s). They had to indicate, using a 7-point Likert scale, how strongly they agree (or disagree) with various statements, e.g. “I was proud of how well the test went”, or “I was angry about the task I had to do” (see Appendix J for the entire list of questions and the instructions).

4 Results


The experiment took place in August 2020 in the BonnEconLab at the University of Bonn.¹⁴ I conducted 52 sessions, with 1 to 6 participants in each session. I collected data from 167 participants in the treatment condition and 55 participants in the control condition. The experiment lasted around 80 minutes and the participants earned 21.25 Euro on average. In the following section, I report the analysis based on the data from 209 participants who correctly answered at least half of the control questions (I excluded 13 participants, that is 5.8% of the sample).

4.1 IQ Test Results and Individual Ranks

Figure 5 presents the distribution of the IQ test scores and ranks assigned to the participants based on the test results. The IQ test score distribution is fairly symmetrical (skewness -0.83), with a mean of 5.13 and a standard deviation of 3.73. The average rank is 5.65 with a standard deviation of 2.67. Importantly, there is no significant difference in the average IQ test score or rank assigned to the participants in the treatment and control group (see Appendix A).

¹⁴Due to the Covid-19 pandemic, I followed special procedures to ensure the safety of participants and others involved. The number of participants per session was restricted to 6 to ensure each participant a place in a separate room. Desks, chairs, and computer equipment were disinfected after every session and the rooms were aired before every session for at least half an hour. At the time of the experiment (August 2020), the Covid-19 pandemic was mostly under control in Germany; the lockdown restrictions were eased, allowing restaurants, schools, and public places to open with appropriate safety measures.

Figure 5: IQ Test Results and Individual Ranks.



4.2 Prior Beliefs about Rank

Before the main task, we elicited from every participant his entire belief distribution. I analyze the data in two ways. Firstly, I look at the aggregate belief distribution. Then, I examine individual distributions and report the averages of individual measures (these include mean belief about rank, median and range).

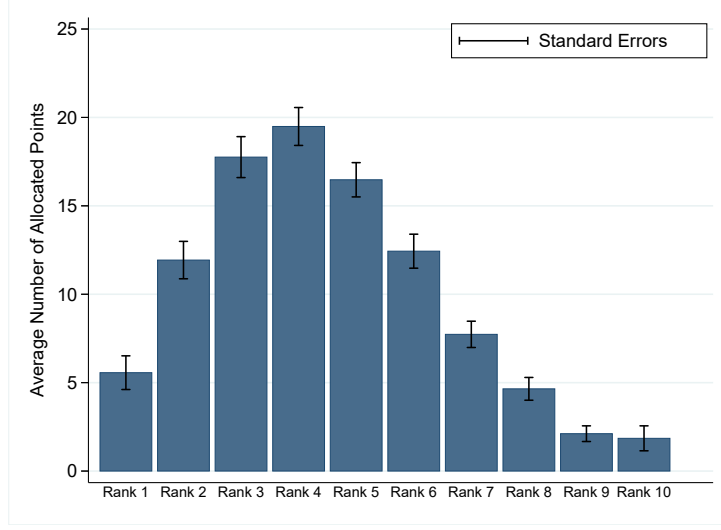
To look at the aggregate of individual belief distributions, I treat separately every decision to allocate x points, $x \in \{0, \dots, 100\}$, to rank k , $k \in \{1, \dots, 10\}$. For each of 10 ranks, I calculate the average number of points allocated by the participants. The resulting aggregate distribution is presented on Panel a) in Figure 6 (each bar indicates the average \pm standard errors). It is visibly skewed to the right, with the mean belief 4.47 and the median 4. On average, the subjects appear to be *overconfident*, as they put a higher probability mass on lower (better) ranks.

In Table 1, I report the averages of individual measures of belief distribution. I look at the average mean belief, median belief, the first and third quartile and range.

Table 1: Individual belief distributions.

	Mean Belief	Q1	Median	Q3	Range
Mean	4.47	3.71	4.45	5.16	4.89
(Std. Dev.)	(1.75)	(1.74)	(1.79)	(1.87)	(1.57)

Figure 6: Average number of allocated points.



Importantly, there is no significant difference between the treatment and the control group (see Appendix A). The averages, however, mask the fact that only 26 participants revealed symmetric belief distribution. Almost half of all subjects (100 participants) revealed positively skewed belief distribution, and the remaining 83 participants revealed negatively skewed belief distribution (the average difference between mean and median in both groups was 0.21).

I define a person to be *overconfident* if his median belief is lower than his true rank. Similarly, I use a term *underconfident* to describe a person who assigns 50% or more probability mass to ranks higher than his true rank. A person is defined to be *unbiased* if his median belief matches his true rank.¹⁵ There is no significant difference in the average bias (defined as a difference between the true rank and the median belief) between the treatment and the control group (see Appendix A).

Using this definition, there are 127 overconfident, 58 underconfident and 24 unbiased participants in my sample. In Appendix G, I report the average rank and measures of individual belief distribution separately for the three types and address the question of apparent overconfidence (Benoît and Dubra, 2011).

¹⁵In common language, Rank 1 denotes “the highest” rank, while Rank 10 is “the lowest”. To avoid confusion, I will not use the customary phrases, but the terms that match the values (for example, a subject whose rank is 5 and median belief is 4 puts higher probability on *lower* ranks).

4.3 Decisions in the Main Task

The main experimental task, neutrally framed as “the second task”, differed depending on the condition. In the treatment condition, subjects observed one number and reported their beliefs about the box from which the number was drawn. In the control condition, participants saw, in random order, numbers from 1 to 10, and stated a report for each one of them. In this subsection, I describe the raw data on received signals and subjects’ reports, as well as the results of the data analysis.

4.3.1 Received Signals

In Figure 7, I present the numbers displayed to participants in the treatment condition depending on the subject’s rank (Panel a) and median belief (Panel b). The size of the hollow circles is proportional to frequencies. Numbers inside the circles denote frequencies. Since more than 50% of participants in the treatment condition saw a number that was equal to their actual rank, circles on the diagonal of Panel a) tend to be larger.

I define a “good” signal as a number lower or equal to one’s median belief. 69 participants (43% of all participants in the treatment condition) received a “good” signal and 91 participants (57% of all) observed a “bad” signal – a number strictly higher than their median belief. The incidence of “good” and “bad” signals in the two conditions is presented in Table 2. Of those who received a “good” signal in the treatment condition,

Figure 7: Signals received in the treatment condition (circle size reflects frequency).

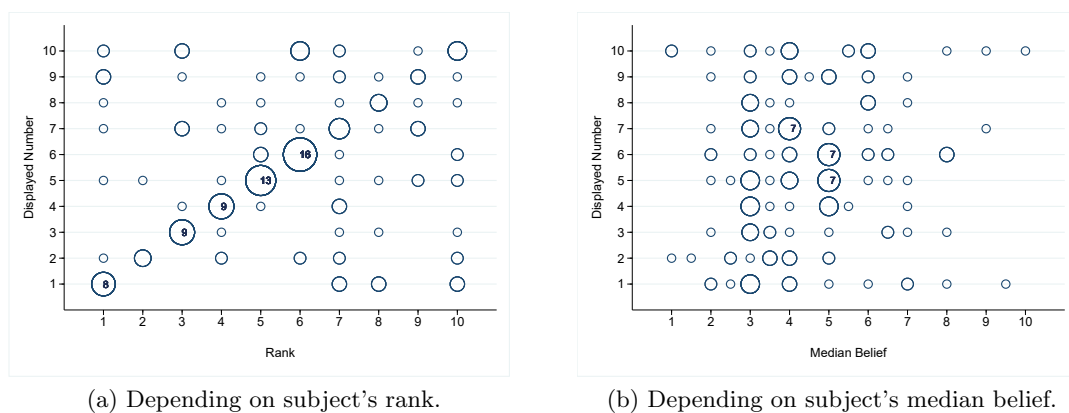


Table 2: Frequency of signals in Treatment and Control.

	Treatment		Control	
	Freq.	Percent	Freq.	Percent
“Good” signals	69	43.13	226	46.12
“Bad” signals	91	56.88	264	53.88
Observations	160		490	

37.7% were classified as overconfident, 46.4% as underconfident, and 15.9% as unbiased. Among participants who obtained a “bad” signal, I classified 78% as overconfident, 14.3% as underconfident, and 7.7% as unbiased.

In the control condition, 226 decisions concerned “good” signals and 264 decisions “bad” ones. Out of 226 decisions regarding “good” signals, 55.8% were made by overconfident agents, 32.3% by underconfident, and 11.9% by unbiased agents. In the case of “bad” signals, 65.9% of decisions were made by overconfident agents, 21.6% by underconfident, and 12.5% by unbiased agents.

Note that, among participants who received a “good” signal, underconfident subjects are *overrepresented* in the treatment condition compared to the control (46.4% versus 32.3%). Among those who received a “bad” signal, overconfident subjects are *overrepresented* in the treatment condition (78% versus 55.8%). If there are differences in how overconfident and underconfident subjects respond to good and bad signals in the two conditions, a simple mean comparison would not recover the true effect of a “good” or “bad” signal on updating. I address this issue in the following sections.

4.3.2 Subjects’ Reports (Raw Data)

In this section, I describe the raw data on the decisions made by participants in the second task. This was our main task: allocating points to Box 1 (with numbers from 1 to 10) and Box 2 (indicating one’s rank) in a way that corresponds to one’s beliefs about the source of the signal. I interpret points allocated to Box 2 as the probability that a subject assigns to the event that the number displayed on the computer screen is his rank. There is a significant difference in the average probability reported in the

Table 3: The average report in the two conditions.

	Mean	Std. Err.	[95% Conf. Int.]		N
Control	29.27	1.42	26.47	32.06	490
Treatment	37.81	2.63	32.66	42.97	160

Note: Standard errors clustered at the participant level.

treatment and the control condition. Participants in the treatment condition reported, on average, 37.81% probability that the signal they received is their rank, while in the control condition the average report was equal to 29.27% (see Table 3).

Figure 8 presents the average number of points allocated to Box 2 after a signal received in the treatment (Panel a) or the control condition (Panel b). The numbers above the x-axis indicate how many participants received a given signal and stated a report. For example, 14 participants in the treatment condition saw “4” displayed on their computer screens and allocated, on average, 65 points to Box 2 (revealing the average subjective probability of 65% that the number “4” is their rank). It is useful to contrast these decisions with the Bayesian benchmark. For each participant, I calculated a Bayesian posterior about the box given his priors and signal realization. The average deviations from the Bayesian update in the two conditions are presented in Figure 9. On Panel c) and d), I separately plotted cases in which subjects assigned zero prior probability to the number displayed on the screen (I refer to them as “outside priors”).

Figure 8: Beliefs about the signal informativeness in the two conditions.

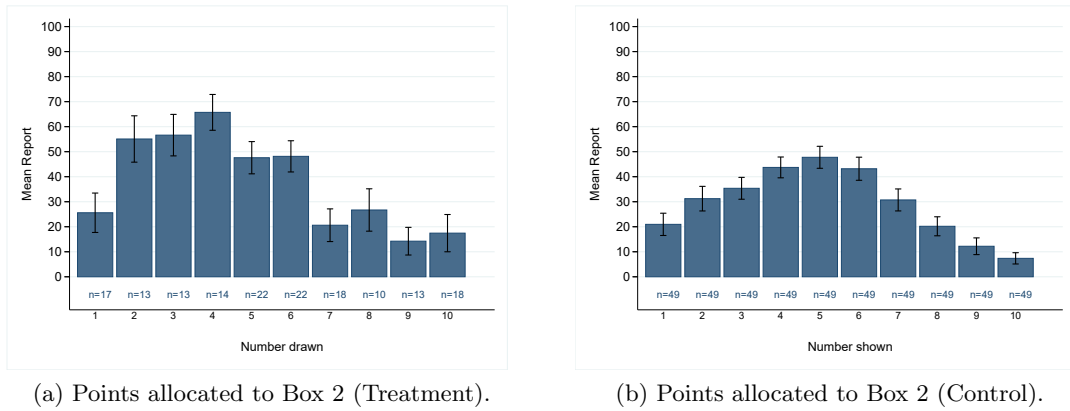
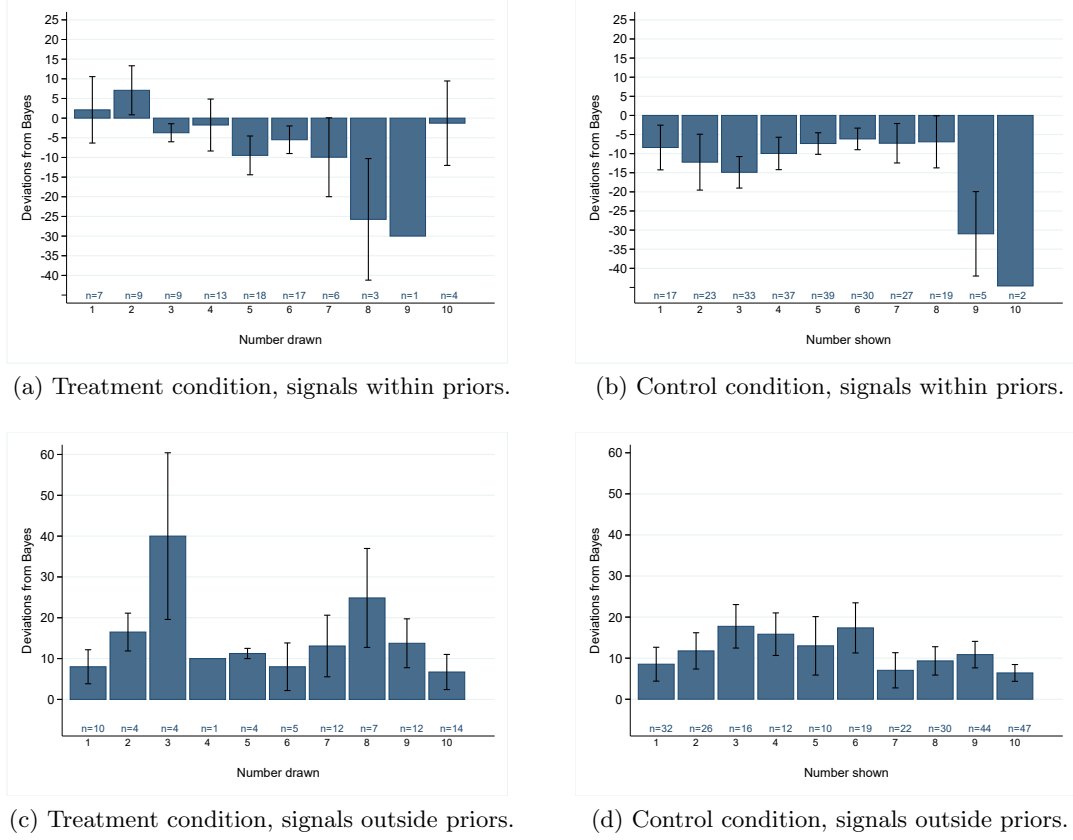


Figure 9: Mean deviation from Bayes for different signals.



The top graphs in Figure 9 give us a glimpse of what is going on. In the control condition, the averages are consistently below zero – subjects tend to allocate fewer points than prescribed by the Bayes’ rule. In the treatment condition, this is true only for higher (worse) signals. In the case of better signals, subjects’ decisions are closer to the Bayesian benchmark. Participants seem to be under-reacting to new information (which may be indicative of conservatism), except when they receive a “good” signal. However, the raw data do not take into account that subjects in the treatment condition were more likely to observe their actual rank. I address this issue in the following section.

4.4 Data Analysis

In this section, I answer my research question: do subjects perceive “good” signals as more informative? I aim at isolating the effect of receiving a “good” signal on subject’s belief about its informativeness revealed through subjects’ choices in the main task.

4.4.1 Matching Estimator

Note that the research question does not refer to the treatment effect itself, but rather the heterogeneity in the treatment effects. Although the assignment into the treatment and control condition is random, the assignment of signals to agents is not. Participants in the treatment condition are presented one number which, with probability $\frac{1}{2}$, is their true rank. This is visible in Panel a) in Figure 7. The hollow circles are much larger on diagonal, meaning that participants are more likely to observe their true rank than any other number.

Imagine a subject who believes his rank is 1. In the control condition, he would consider all ten numbers, and 9 out of 10 decisions would pertain to an unfavorable signal. On the other hand, if he was in the treatment condition and his rank was indeed 1, he would receive a bad signal with much lower probability: $\frac{1}{2} \times \frac{9}{10}$. This leads to the covariance between the treatment status and signals considered by the participants. If there are reasons to believe that people with different beliefs or ranks respond differently to good signals, a simple comparison of means would not recover the treatment effect. I present this argument formally in Appendix D.1.

Moreover, the mapping from prior belief distribution and rank to belief about the box is likely to be non-linear. As a consequence, the OLS estimates may not be efficient. Nevertheless, I conduct regression analysis and report the results in Appendix C.

For these reasons, I use a different approach to analyze the data. I follow Heckman et al. (1998) and construct a matching estimator:

$$\hat{Y}_i^N = \sum_{j=1}^J w_j^i Y_j^C, \quad (1)$$

where \hat{Y}_i^N denotes beliefs of subject i from the treatment condition if he had not received the signal (the counterfactual outcome), Y_j^C denotes beliefs of subject j in the control condition (it includes a correction for potential bias as in Abadie and Imbens, 2011), $j = \{1, \dots, J\}$, and w_j^i is the weight assigned to j in the counterfactual outcome of subject i . The weights are normalized such that $0 \leq w_j \leq 1$ and $\sum_{j=1}^J w_j = 1$. I

estimate the weights using a kernel regression for each participant in the treatment condition. I describe the estimation procedure in detail in Appendix D.2.

Intuitively, I construct the counterfactual to the participant i in the following way: I take the decisions of *all* participants in the control condition regarding the number that the participant i saw. However, not all observations in the control condition receive the same weight. Those participants whose true ranks and prior beliefs were closer to that of the participant i , receive a higher weight.¹⁶ I interpret the counterfactual as what would subject i decide if he was in the control condition. Having constructed this counterfactual scenario, I look at the effect of a “good” signal on updating using regression analysis.

4.4.2 Regressions Analysis

The results of regression analysis using counterfactual outcomes are reported in Table 4. The dependent variable is the difference in points allocated to Box 2 (indicative of one’s rank) in the treatment and the counterfactual scenario. I interpret it as the difference in reported probabilities that a signal is entirely informative about one’s rank after receiving it, compared to what they would conclude if they considered the same signal in the control condition.

In the first specification, I regress this difference on a constant. The coefficient is significant and is equal to 4.95, a value similar to the one obtained in the regression based on all observations from the control condition. Subjects reported around 5 percentage points higher probability that the signal is their rank after receiving it. However, the effect is entirely driven by the response to “good” signals. In the second specification, I add an indicator variable “Good Signal” which takes the value of 1 if a signal received was lower or equal to the subject’s median belief. The coefficient at the Good Signal is positive and significant. After receiving a “good” signal, participants tend to put 10.55 higher probability on the signal being their rank in the treatment condition compared to what they would decide ex ante. There is no significant difference after “bad” signals.

¹⁶In the baseline specification, I match subjects using their true rank and prior beliefs. In Appendix F, I report the results based on two alternative specifications: in Specification 2, I match participants based on their true rank and prior beliefs about the number under consideration, and in Specification 3, I use only the distribution of prior beliefs (in theory, it subsumes information that a subject has about his performance). The results are very similar to the baseline specification.

Table 4: The effect of the signal’s valence.

	(1)	(2)	(3)	(4)
Good Signal		10.55*** (3.87)	7.72* (3.94)	15.68*** (5.05)
Outside Priors			-10.59*** (3.92)	-2.91 (4.96)
Outside Priors \times Good				-19.40** (7.89)
Constant	4.95** (1.96)	0.40 (2.54)	6.45* (3.35)	2.06 (3.75)
Observations	160	160	160	160

Standard errors in parentheses

* $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$

Note: The dependent variable is the difference between numbers of points allocated to Box 2 in the treatment and in the counterfactual (kernel-based matching). “Good Signal” indicator variable takes value 1 if the signal was below or equal to the median of subject’s belief distribution, and 0 otherwise. “Outside priors” indicator variable takes value 1 if the subject attached a zero prior probability to the signal being his rank, and 0 otherwise.

In the third specification, I control for signals that were outside of subjects’ prior belief distribution. I add a dummy variable “Outside Priors” taking the value of 1 if a subject assigned a prior probability of zero to the signal displayed on his screen. However, a signal outside of priors can also be good or bad and can have a different effect depending on its valence. In the last specification, I add an interaction term of the Good Signal and the Outside Prior variable. The estimated effect is negative and counteracts the positive effect of the Good Signal variable. The coefficient at the Good Signal variable is higher compared to previous estimates. If subjects assigned a non-zero prior probability to the “good” signal displayed on their computer screen, they reported a 15-percentage-point higher probability that the signal was their rank in the treatment condition compared to what they would decide in the control condition. If the signal was outside the subject’s prior belief distribution, the effect is entirely reduced.¹⁷ There is no significant difference in the case of “bad” signals.

¹⁷The results suggest that the good-news effect is not universal across signals. Getting a signal “too good to be true” makes the agent skeptical and leads him to assign a lower probability than he would in the control condition. It does not contradict the theory, as for the subject to experience the belief-based utility it is necessary that the signal affects his beliefs, and it may not be the case if the signal is outside of the subject’s priors.

The standard errors reported in Table 4 do not account for the matching procedure. There is additional uncertainty coming from that we do not know how well our control group reproduce the counterfactual outcome. Abadie and Imbens (2006) derive analytical formulas for a consistent estimator for the large-sample variance of the nearest-neighbor matching estimator. However, large-sample techniques may not be well suited when the number of units in the comparison group is small (Abadie et al., 2010).¹⁸ For this reason, I employ an inferential technique proposed in Abadie et al. (2010) that I describe in the following section.

4.4.3 Placebo Studies

For a robust inference in a finite sample, Abadie et al. (2010) propose an inferential technique based on “placebo studies”. The idea behind it is to compare the actual treatment effect to the distribution of so-called “placebo” treatment effects. The latter is calculated by assigning the treatment status to a random sample of all observations, conducting the same analysis and storing the estimated coefficients. I provide the details of the procedure in Appendix D.3. I focus on Specification 4) from the previous section.

Figures 10 and 11 summarize the results of the placebo studies. In Figure 10, I present a histogram of coefficients at the Good Signal variable. Figure 11 shows the coefficients at the interaction of the Good Signal and the Outside Prior variable. The vertical lines denote the actual treatment effects. One can notice that their magnitudes are extreme relative to the distributions of coefficients in the placebo studies, indicating the statistical significance of the actual treatment effects. The empirical distribution of the placebo effects allows me to calculate the p-value of a two-sided test to assess the statistical significance of the actual treatment effect. Formally, I test a hypothesis that there is no difference between the actual treatment effect and the placebo treatment effect. The corresponding p-values are 0.003 in the case of the Good Signal variable and 0.039 for the interaction term. I conclude that both effects are statistically significant.

¹⁸Although the sample size of 209 subjects would not be considered small for an experimental economist, it is a small sample given our set-up. If we divide participants based on their prior belief distribution, observed signal, and its relation to the subject’s priors, we end up with much smaller comparison groups.

Figure 10: Distribution of coefficients at the Good Signal variable (Specification 4).





Figure 11: Distribution of coefficients at the interaction term (Specification 4).




4.5 Manipulation Check

I argue that the treatment effect is caused by the utility from beliefs induced by the signal. First of all, I provide evidence that the signal received in the treatment condition affected subjects' beliefs. In Figure 12, I present subjects' beliefs before and after the main task, that is, beliefs revealed in the first and the second belief elicitation. The graph shows points allocated to the rank corresponding to the number displayed on subjects' screens. I compare these values to the counterfactual: how many points they would allocate to the respective ranks if they did not receive a signal. There is no change in beliefs in the counterfactual scenario (denoted with a blue line). In the treatment condition, the change in beliefs is significant (marked in red on the graph). Subjects allocated almost two times as many points in the second belief elicitation to the rank displayed on the screen. In Appendix B, I separately plotted changes in beliefs after “good” and “bad” signals, and signals that were within subjects' prior belief distributions.

Secondly, I exclude alternative hypotheses. One may worry that subjects in the control condition exerted less effort per decision (e.g. due to increasing marginal cost of effort or lower monetary incentives in the control condition). To alleviate this concern, we asked participants in the treatment condition, before they received an actual signal,

Figure 12: Beliefs before and after the signal.



to consider about every possible signal realization. We showed them, one by one, every possible number and asked them to think what they would do if this number was drawn later. This additional part makes the total time spent on the second task similar in both conditions (see Table 9 in Appendix B).

One may argue that the total time spent on the task may not be a perfect measure of effort and there still may be differences in cognitive effort exerted when making a decision in the treatment and in the control condition. However, if this was the case, one would expect larger deviations from the rational benchmark in the control condition. As reported in Table 10 in Appendix B, there is no significant difference in absolute deviations from the Bayesian benchmark in the two conditions. I provide additional evidence to support my interpretation of the results as being driven by changes in belief-based utility in Section 5.

4.6 Earnings

In this section, I look at the payoffs from the main task in the treatment and the control conditions. In both conditions, subjects were remunerated with “lottery tickets”: a higher probability of receiving a large reward of 12 Euro. Firstly, I compare the average probability obtained in the treatment and in the control condition, as well as in the treatment and in the counterfactual scenario. Participants in the treatment condition obtained, on average, 65.5% probability of receiving a large reward, a probability that was 12.6 percentage points lower (19.2% in relative terms) than the probability earned by the subjects in the control condition ($p\text{-value} = 0.000$). If we compare the treated subjects to the control constructed in Section 4, participants in the treatment condition earned 8.7 percentage points less (13.3% in relative terms) than they would have in the counterfactual ($p\text{-value} = 0.006$).

One can conclude that, on average, subjects in the treatment condition are worse-off (in monetary terms) than they would be if they were not given an opportunity to observe a signal and learn. The result gives rise to several questions: to what extent subjects are aware of their propensity to interpret realized signals differently? If possible, would they commit to their decisions *ex ante*? The result and its implications suggest interesting directions for future research.

5 Additional Evidence

In this section, I examine a complementary data set of subjects’ answers to questionnaires described in Section 2. Firstly, I look at the subjects’ personality traits, anxiety levels, as well as habitual use of emotion regulation strategies, and report their correlations with subjects’ decisions in the second task.

5.1 Emotion Regulation Questionnaire

In this section, I examine subjects’ answers to the emotion regulation questionnaire, BIG-5 and STAI. In Table 5, I report correlations between subjects’ decisions in the treatment condition (relative to the Bayesian benchmark) and the above-mentioned measures. The absolute deviations from Bayesian updating are correlated with the habitual use of reappraisal. The coefficient value of -0.18 indicates a weak, negative correlation significant at the 0.05 level.

Table 5: Deviations from rationality and agents’ characteristics in the treatment condition.

	DevB	Extr	Cons	Open	Neur	Agre	Trait	State	Reapp	Supr
DevB	1.00									
Extr	0.00	1.00								
Cons	0.05	-0.01	1.00							
Open	-0.09	0.22*	0.10	1.00						
Neur	0.12	-0.24*	-0.26*	0.16*	1.00					
Agre	-0.03	0.07	0.07	0.07	-0.13	1.00				
Trait	-0.07	0.29*	0.35*	-0.09	-0.71*	0.23*	1.00			
State	-0.15	0.28*	0.17*	-0.03	-0.58*	0.24*	0.70*	1.00		
Reapp	-0.18*	0.09	0.15	0.18*	-0.17*	0.22*	0.13	0.17*	1.00	
Supr	-0.04	-0.19*	0.05	-0.17*	-0.04	0.03	-0.13	-0.14	0.38*	1.00

* $p < 0.05$

Note: “DevB” stands for deviations from Bayesian update. I use the labels: “Extr”, “Cons”, “Open”, “Neur”, and “Agre” for BIG-5 personality traits: extraversion, conscientiousness, openness to experience, agreeableness and neuroticism, respectively. I denote Anxiety trait and state with “Trait” and “State” (the two measures are defined such that a higher score indicates less anxious individual). “Reapp” and “Supr” stands for emotion regulation strategies: reappraisal and suppression.

In Table 6, I present the estimates of regressions based on decisions made by participants in the treatment condition. I regress the independent variable, the absolute deviations from Bayesian update, on the independent variable “Reappraisal” that mea-

sures subject’s habitual use of reappraisal. The coefficient at the Reappraisal variable is negative and significant at the 0.05 level. Reporting one point higher response on the 7-point Likert scale in questions about one’s habitual use of reappraisal leads to a 3-point decrease in the distance from Bayesian update. The value doesn’t change much if I control for subject’s rank, median belief or whether the signal he received was below or above his median belief or not within the prior belief distribution. The results show that subjects’ decisions correlate with the way they handle positive and negative emotions in their daily life. The more used they are to regulate their emotions by thinking differently about the situation they found themselves in, the more they adhere to rational decision-making. To investigate this further, I take a closer look at emotion regulation strategies together with self-reported emotions experienced before the task.

Table 6: The effect of reappraisal on deviations from Bayes.

	(1)	(2)
Reappraisal	-2.96** (1.29)	-2.82** (1.29)
Constant	26.61*** (5.76)	27.33*** (7.50)
Controls	No	Yes
Observations	160	160

Standard errors in parentheses

* $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$

Note: The dependent variable is absolute deviations from the Bayesian update. Controls include the subject’s rank and median prior belief, a dummy variable equal 1 if the signal was below or equal to the median prior belief, and a dummy variable equal 1 if the signal was outside of the subject’s prior beliefs.

5.2 Test-related Emotions

In addition to the data presented so far, I collected survey data about test-related emotions experienced by participants before receiving the signal.¹⁹

¹⁹In the instructions displayed on the screen, I highlighted that questions refer to the particular moment in time: *after* learning the nature of the task, but *before* seeing the number.

Out of eight test-related emotions, anxiety and hopelessness significantly correlate with absolute deviations from Bayesian updating in the treatment condition. However, when I regressed absolute deviations from Bayesian updating on all test-related emotions, only hopelessness was highly statistically significant ($p\text{-value} = 0.02$) and remained so, even after adding additional controls on subjects' rank, median belief, and signal's value or its relation to the subject's beliefs.

Hopelessness was measured by agreement with the statement "I felt that I would rather not do this part because I've lost all hope.". As reported in the first column in Table 7, stating a 1-point higher answer to the question translates to an increase of 4.3 points in absolute deviation from Bayesian updating (controlling for all remaining test-related emotions). The coefficient at the Hopelessness variable remains unchanged if I control for the emotion regulation strategies: suppression and reappraisal (Specification 2) in Table 7. Of the two strategies, only reappraisal is different from zero and significant.

Table 7: The effect of self-reported emotions on deviations from rationality.

	(1)	(2)	(3)
Hopelessness	4.31** (1.83)	4.30** (1.82)	17.23*** (4.62)
Reappraisal		-2.82** (1.42)	2.21 (2.16)
Hopelessness \times Reappraisal			-3.10*** (1.02)
Constant	10.00 (8.28)	20.18** (10.11)	2.73 (11.40)
Controls 1	Yes	Yes	Yes
Controls 2	No	Yes	Yes
Observations	160	160	160

Standard errors in parentheses

* $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$

Note: The dependent variable is absolute deviations from the Bayesian update. The independent variable "Hopelessness" was measured by the extent to which a subject agreed with the statement "I felt that I would rather not do this part because I've lost all hope.". "Reappraisal" refers to self-reported habitual use of reappraisal. Controls 1 include all other emotions reported by subjects; Controls 2 include the measure of habitual use of suppression.

Moreover, it has the expected negative sign and value similar to that reported in Table 5. I hypothesize that the use of reappraisal counteracts the negative impact of Hopelessness. To test this hypothesis, I add to the regression the interaction of Hopelessness and Reappraisal. I report the estimation results in the last column of Table 7. The coefficient at the interaction term is negative and highly significant, whereas the coefficient at the Reappraisal variable loses its significance. At the same time, the coefficient at Hopelessness increases fourfold and gains significance, suggesting that its impact is much larger without the offsetting effect of reappraisal.

While only suggestive, the evidence presented in this section supports the view that the treatment effect is stemming from the visceral, emotion-based reaction to signals. That reaction lies at the heart of what economists call “the belief-based utility” and is the driving force behind asymmetric updating.

6 Conclusions

In this paper, I propose a new test of the hypothesis that people interpret positive feedback as more informative. To this end, I designed a simple experiment with two conditions. In the treatment condition, participants observe a signal about their intelligence and decide whether the signal is informative or not. In the control condition, participants make the same choice without receiving a factual signal: they are asked to specify their actions conditioning on possible signal realizations. This design allows me not only to pin down the causal effect of signal’s valence on updating but also to uncover the underlying mechanism.

The experimental data reveal that people tend to interpret favorable signals as more informative due to the changes in belief-based utility. Participants reported a 27% higher probability of a positive signal being entirely informative about their rank after receiving it, compared to what they would conclude *ex ante*, without observing its realization. As a result of distorted perception of positive signals, subjects in the treatment condition ended up with *lower* payoffs (by 13%) compared to what they would get if they were in the control condition. A striking conclusion of the experiment is that allowing subjects to acquire signals and learn can make them worse-off (in monetary terms).

There is mounting evidence that people derive utility not only from physical outcomes but also from their beliefs about the current or future state. The belief-based utility is likely to be a driving force behind overconfidence, the demand for (and avoidance of) information, and belief polarization. Yet, the way it influences people’s actions and beliefs is not fully understood. My study takes the next step towards explaining its role by revealing how belief-based utility shapes the way we interpret new information.

References

- Abadie, Alberto, Alexis Diamond, and Jens Hainmueller (2010). “Synthetic control methods for comparative case studies: Estimating the effect of California’s tobacco control program”. In: *Journal of the American statistical Association* 105.490, pp. 493–505.
- Abadie, Alberto and Guido W. Imbens (2006). “Large sample properties of matching estimators for average treatment effects”. In: *econometrica* 74.1, pp. 235–267.
- (2011). “Bias-corrected matching estimators for average treatment effects”. In: *Journal of Business & Economic Statistics* 29.1, pp. 1–11.
- Ambuehl, Sandro and Shengwu Li (2018). “Belief updating and the demand for information”. In: *Games and Economic Behavior* 109, pp. 21–39.
- Barber, Brad M and Terrance Odean (2001). “Boys will be boys: Gender, overconfidence, and common stock investment”. In: *The quarterly journal of economics* 116.1, pp. 261–292.
- Bénabou, Roland and Jean Tirole (2016). “Mindful Economics: The Production, Consumption, and Value of Beliefs”. In: *Journal of Economic Perspectives* 30.3, pp. 141–164.
- Benoît, Jean-Pierre and Juan Dubra (2011). “Apparent overconfidence”. In: *Econometrica* 79.5, pp. 1591–1625.
- Brandts, Jordi and Gary Charness (2009). *The strategy method: A survey of experimental evidence*. Tech. rep. mimeo, Department of Business Economics U. Autònoma de Barcelona and ...
- Brunnermeier, Markus K and Jonathan A Parker (2005). “Optimal expectations”. In: *American Economic Review* 95.4, pp. 1092–1118.
- Buser, Thomas, Leonie Gerhards, and Joël Van Der Weele (2018). “Responsiveness to feedback as a personal trait”. In: *Journal of Risk and Uncertainty* 56.2, pp. 165–192.
- Camerer, Colin and Dan Lovallo (1999). “Overconfidence and Excess Entry: An Experimental Approach”. In: *The American Economic Review* 89.1, pp. 306–318.

- Caplin, Andrew and John V Leahy (2019). *Wishful Thinking*. Tech. rep. National Bureau of Economic Research.
- Chew, Soo Hong, Wei Huang, and Xiaojian Zhao (2019). “Motivated false memory”. In: *Available at SSRN 2127795*.
- Coutts, Alexander (2019). “Good news and bad news are still news: Experimental evidence on belief updating”. In: *Experimental Economics* 22.2, pp. 369–395.
- Coutts, Alexander, Leonie Gerhards, and Zahra Murad (2020). “What to blame? Self-serving attribution bias with multi-dimensional uncertainty”. In:
- Eil, David and Justin M. Rao (2011). “The good news-bad news effect: Asymmetric processing of objective information about yourself”. In: *American Economic Journal: Microeconomics* 3.2, pp. 114–138.
- Engelmann, Jan, Maël Lebreton, Peter Schwardmann, Joel J van der Weele, and Li-Ang Chang (2019). “Anticipatory anxiety and wishful thinking”. In:
- Enke, Benjamin and Thomas Graeber (2019). *Cognitive uncertainty*. Tech. rep. National Bureau of Economic Research.
- Ertac, Seda (2011). “Does self-relevance affect information processing? Experimental evidence on the response to performance and non-performance feedback”. In: *Journal of Economic Behavior & Organization* 80.3, pp. 532–545.
- Gerlitz, Jean-Yves and Jürgen Schupp (2005). “Zur Erhebung der Big-Five-basierten persönlichkeitsmerkmale im SOEP”. In: *DIW Research Notes* 4, p. 2005.
- Gilboa, Itzhak and Massimo Marinacci (2016). “Ambiguity and the Bayesian paradigm”. In: *Readings in formal epistemology*. Springer, pp. 385–439.
- Golman, Russell, David Hagmann, and George Loewenstein (2017). “Information avoidance”. In: *Journal of Economic Literature* 55.1, pp. 96–135.
- Gross, James J and Oliver P John (2003). “Individual differences in two emotion regulation processes: implications for affect, relationships, and well-being.” In: *Journal of personality and social psychology* 85.2, p. 348.
- Grossman, Zachary and David Owens (2012). “An unlucky feeling: Overconfidence and noisy feedback”. In: *Journal of Economic Behavior & Organization* 84.2, pp. 510–524.
- Hahn, Jinyong (1998). “On the role of the propensity score in efficient semiparametric estimation of average treatment effects”. In: *Econometrica*, pp. 315–331.
- Hastie, Trevor, Robert Tibshirani, and Jerome Friedman (2009). *The elements of statistical learning: data mining, inference, and prediction*. Springer Science & Business Media.
- Heckman, James J, Hidehiko Ichimura, and Petra Todd (1998). “Matching as an econometric evaluation estimator”. In: *The review of economic studies* 65.2, pp. 261–294.

- Hestermann, Nina and Yves Le Yaouanq (2020). “Experimentation with Self-Serving Attribution Biases”. In: *American Economic Journal: Microeconomics*.
- Hossain, Tanjim and Ryo Okui (2013). “The binarized scoring rule”. In: *Review of Economic Studies* 80.3, pp. 984–1001.
- Huffman, David, Collin Raymond, and Julia Shvets (2019). “Persistent overconfidence and biased memory: Evidence from managers”. In: *Pittsburgh: University of Pittsburgh*.
- Kőszegi, Botond (2006). “Ego Utility, Overconfidence, and Task Choice”. In: *Journal of the European Economic Association* 4.June, pp. 673–707.
- Lerner, Jennifer S, Ye Li, Piercarlo Valdesolo, and Karim S Kassam (2015). “Emotion and decision making”. In: *Annual review of psychology* 66.
- Malmendier, Ulrike and Geoffrey Tate (2005). “Does overconfidence affect corporate investment? CEO overconfidence measures revisited”. In: *European Financial Management* 11.5, pp. 649–659.
- (2008). “Who makes acquisitions? CEO overconfidence and the market’s reaction”. In: *Journal of financial Economics* 89.1, pp. 20–43.
- Mezulis, Amy H, Lyn Y Abramson, Janet S Hyde, and Benjamin L Hankin (2004). “Is there a universal positivity bias in attributions? A meta-analytic review of individual, developmental, and cultural differences in the self-serving attributional bias.” In: *Psychological bulletin* 130.5, p. 711.
- Möbius, Markus M, Muriel Niederle, Paul Niehaus, and Tanja S Rosenblat (2014). “Managing Self-Confidence”. Working Paper.
- Niederle, Muriel and Lise Vesterlund (2007). “Do women shy away from competition? Do men compete too much?” In: *The Quarterly Journal of Economics* 122.3, pp. 1067–1101.
- Ortoleva, Pietro and Erik Snowberg (2015). “Overconfidence in political behavior”. In: *American Economic Review* 105.2, pp. 504–35.
- Pekrun, Reinhard, Thomas Goetz, Anne C Frenzel, Petra Barchfeld, and Raymond P Perry (2011). “Measuring emotions in students’ learning and performance: The Achievement Emotions Questionnaire (AEQ)”. In: *Contemporary educational psychology* 36.1, pp. 36–48.
- Rosenbaum, Paul R and Donald B Rubin (1983). “The central role of the propensity score in observational studies for causal effects”. In: *Biometrika* 70.1, pp. 41–55.
- Schwardmann, Peter and Joel Van der Weele (2019). “Deception and self-deception”. In: *Nature human behaviour* 3.10, pp. 1055–1061.
- Spielberger, Charles D (1983). “State-trait anxiety inventory for adults”. In:
- Uusberg, Andero, Jamie L Taxer, Jennifer Yih, Helen Uusberg, and James J Gross (2019). “Reappraising reappraisal”. In: *Emotion Review* 11.4, pp. 267–282.

- Van den Steen, Eric (2004). “Rational overoptimism (and other biases)”. In: *American Economic Review* 94.4, pp. 1141–1151.
- Zimmermann, Florian (2020). “The dynamics of motivated beliefs”. In: *American Economic Review* 110.2, pp. 337–61.

A Differences between the treatment and the control group

Differences between participants in Treatment and Control.

	Treatment	Control	p-value		
			H_0 : Diff < 0	Diff \neq 0	Diff > 0
IQ score	5.12 (0.30)	5.16 (0.50)	0.47	0.94	0.53
Rank	5.59 (0.21)	5.82 (0.39)	0.31	0.61	0.69
Bias	1.18 (0.22)	1.23 (0.43)	0.46	0.91	0.54
Absolute Bias	2.38 (0.14)	2.60 (0.28)	0.24	0.47	0.76
N	160	49			

Note: “Bias” is defined as difference between rank and median belief. Standard errors in parenthesis.

Table 8: Measures of Individual Prior Belief Distributions.

	Treatment	Control	H_0 : Diff < 0	p-value	
				Diff \neq 0	Diff > 0
Prior Beliefs:					
Mean	4.43 (0.14)	4.56 (0.26)	0.33	0.65	0.67
1 st Quartile	3.69 (0.13)	3.79 (0.27)	0.35	0.70	0.65
Median	4.41 (0.13)	4.58 (0.27)	0.28	0.56	0.72
3 rd Quartile	5.11 (0.15)	5.34 (0.27)	0.23	0.45	0.77
N	160	49			

B Manipulation Check

Table 9: Decision time and time spent on the main task, in seconds.

	Treatment	Control	p-value		
			H_0 : Diff < 0	Diff \neq 0	Diff > 0
Total time	249.42 (10.66)	228.56 (17.64)	0.83	0.34	0.17
Total time (corrected)	231.89 (10.50)	228.56 (17.64)	0.56	0.88	0.44
N	160	49			
Decision time	40.49 (6.63)	22.86 (1.13)	0.995	0.01	0.005
Decision time (matched data)	40.49 (6.63)	24.21 (0.57)	0.993	0.015	0.007
N	160	490			

Table 10: Deviations from Bayes in the main task (Treatment vs Control).


<i>Dependent variable: absolute difference between subjects' reports and the Bayesian benchmark.</i>	
	(1)
Treatment	-0.46 (1.72)
Constant	14.23*** (0.93)
Observations	650

Standard errors clustered at the participant level.

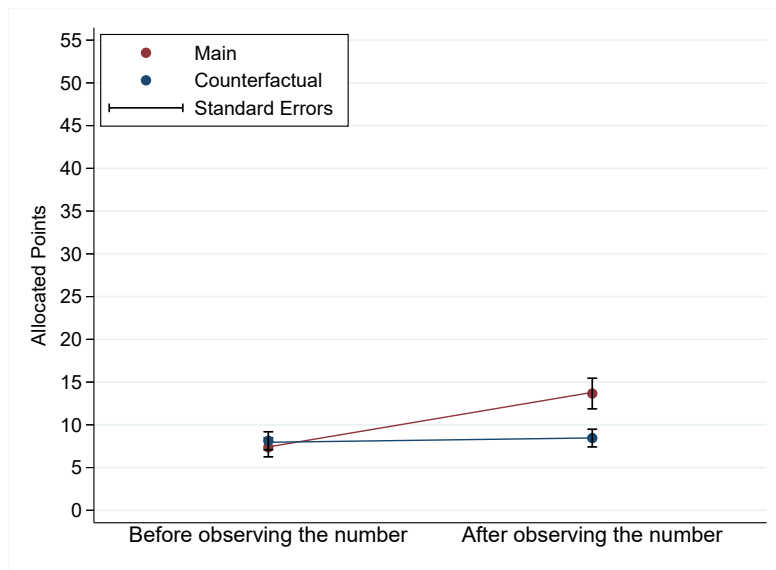
Their values in parentheses.

* $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$

Figure 13: Change in beliefs measured before and after the main task.




(a) After “good” signals.




(b) After “bad” signals.

Figure 14: Change in beliefs after signals “within priors”.



(a) After “good” signals within priors.



(b) After “bad” signals within priors.

C Regression Analysis

In this section, I describe the regression analysis of subjects' decisions in the second task. The dependent variable is the number of points allocated to Box 2, which reflects the probability that a subject assigns to the signal being his rank. Since the decision depends on the subject's prior beliefs about his rank, I regress it on Bayesian updated belief about the box that accounts for subjects' priors.

In Table 11, I present the estimation results of different specifications. Firstly, I regress the dependent variable on the Bayesian belief (the independent variable "Bayes") and the treatment dummy (I refer to it as the "Treatment" variable). As reported in the first column, both coefficients are positive and highly significant.

Table 11: The effect of receiving a signal on decisions.

	(1)	(2)	(3)	(4)	(5)	(6)
Bayes	0.69*** (0.04)	0.83*** (0.09)	0.67*** (0.04)	0.79*** (0.09)	0.67*** (0.04)	0.79*** (0.09)
Treatment	4.55** (1.95)	4.45** (1.93)	4.84** (1.91)	4.74** (1.89)	1.62 (2.60)	1.58 (2.59)
Good Signal			5.08** (2.14)	4.82** (2.18)	3.29 (2.59)	3.07 (2.63)
Treatment \times Good					7.36* (4.27)	7.22* (4.27)
Outside Priors		9.69* (5.23)		8.88 (5.38)		8.74 (5.34)
Constant	9.49*** (1.29)	0.56 (5.30)	7.85*** (1.46)	-0.25 (5.34)	8.69*** (1.64)	0.70 (5.40)
Observations	650	650	650	650	650	650

Standard errors clustered at individual level. Their values in parentheses.

* $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$

Note: The dependent variable is the number of points allocated to Box 2. The independent variable "Bayes" refers to the Bayesian prediction based on the prior belief distribution. "Treatment" is a dummy variable indicating the Treatment condition. "Good Signal" dummy variable takes value 1 if the signal was below or equal to the median of the subject's belief distribution and 0 otherwise. "Outside priors" dummy variable indicates that the subject attached a prior probability of 0 to the signal being his rank.

The coefficient at the Bayes variable is smaller than 1 (p-value = 0.000), implying that, on average, subjects allocate fewer points to Box 2 (with numbers equal to their rank) than a Bayesian decision-maker would do. However, subjects in the treatment condition tend to allocate 4.5 points more to Box 2 than in the control condition.

The coefficient at the Bayes variable significantly increases when I add a dummy “Outside Priors” that takes the value of 1 if a signal was outside of the subject’s priors (Specification 2). On average, subjects allocated 9.69 points to Box 2 after a signal to which they assigned a prior probability of zero (the coefficient is significant at 0.1 level).

In the third specification, I add a dummy variable “Good”, which takes the value of 1 if a number displayed on the individual screen was lower or equal to the subject’s median belief – a favorable signal about one’s rank. The coefficient is positive and significant, implying that participants allocated on average more points to the box indicative of their rank after good news. Its value remains unchanged if I control for the signal being outside of the subject’s prior distribution (Specification 4).

Finally, I test whether subjects treated good signals differently in the treatment and control conditions. In Specification 5), I add a variable “Treatment \times Good” that is an interaction of the Treatment and Good Signal variables. As I see in column 5 in Table 11, the coefficient at the interaction term is high and significant at the 0.1 significance level, while the coefficients at the Treatment and Good Signal variables drop and lose their significance.

D Estimation Procedure

In this section, I provide more details on the estimation procedure. I begin by motivating the use of a matching estimator and contrasting it with a linear regression. Next, I describe the procedure that I use to determine the parameters of the matching estimator. Finally, I discuss the inference.

I adopt the standard model of treatment effects (Abadie and Imbens, 2006; Heckman et al., 1998; Rosenbaum and Rubin, 1983). Let $Y_i(W_i)$ denote the outcome of interest: the number of points that a participant i allocated to Box 2. W_i is a binary variable indicating whether the subject was assigned to the treatment ($W_i = 1$) or the control condition ($W_i = 0$). The average treatment effect is defined as

$$\tau = \mathbb{E}[Y_i(1) - Y_i(0)]. \quad (2)$$

Let subjects' decisions be described by an additive model

$$\begin{aligned} Y_i(1) &= \mu_1(\mathbf{X}_i) + \varepsilon_1 \\ Y_i(0) &= \mu_0(\mathbf{X}_i) + \varepsilon_0, \quad \varepsilon_1, \varepsilon_0 \sim i.i.d \end{aligned} \quad (3)$$

where μ_1 and μ_0 are unknown functions of a k -dimensional vector of individual characteristics \mathbf{X} .

Because of a random assignment to the two conditions, the sample average of outcomes recorded in the control condition is an unbiased estimator of the counterfactual outcome $Y_i(0)$. Therefore, a consistent estimation of the treatment effect entails a simple comparison of mean outcomes in both groups of participants

$$\hat{\tau} = \frac{1}{|N_T|} \sum_{i \in N_T} Y_i - \frac{1}{|N_C|} \sum_{i \in N_C} Y_i, \quad (4)$$

where N_T and N_C denote the set of participants in the treatment and the control condition respectively, and $|A|$ is the cardinality of a set A .

However, I am interested in heterogeneous treatment effects, defined as

$$\tau(\mathbf{x}) = \mathbb{E}[Y_i(1) - Y_i(0) | \mathbf{X}_i = \mathbf{x}]. \quad (5)$$

In case of a random assignment, one can simply compare means

$$\tau(\mathbf{x}) = \mathbb{E}[Y_i | W_i = 1, \mathbf{X}_i = \mathbf{x}] - \mathbb{E}[Y_i | W_i = 0, \mathbf{X}_i = \mathbf{x}]. \quad (6)$$

Note that the conditional expectation satisfy $\mathbb{E}[Y_i | W_i = w, \mathbf{X}_i = \mathbf{x}] = \mu_w(\mathbf{x})$.

D.1 Potential Problem: Selection

Estimation of $\mu_w(\mathbf{x}) = \mathbb{E}[Y_i | W_i = w, \mathbf{X}_i = \mathbf{x}]$ turns out to be quite challenging. To illustrate why the OLS estimate may not be consistent, let's consider the task of estimating the treatment effect among participants who received a good signal. Consequently, X_i consists of a dummy variable equal to 1 for people who received a signal above or equal to their median belief, and 0 otherwise. The treatment effects are defined as

$$\tau(1) = \mathbb{E}[Y_i | W_i = 1, X_i = 1] - \mathbb{E}[Y_i | W_i = 0, X_i = 1]$$

and

$$\tau(0) = \mathbb{E}[Y_i | W_i = 1, X_i = 0] - \mathbb{E}[Y_i | W_i = 0, X_i = 0].$$

We are interested in the difference $\tau(1) - \tau(0)$.

Although the assignment to the treatment and control conditions is random, the experimental design may lead to the covariance between the treatment status and signals received by participants. This correlation may arise because participants in the treatment condition are presented one number, which is their rank with probability $\frac{1}{2}$. Therefore, underconfident agents will see a good signal more often than overconfident agents. In contrast, in the control conditions subjects see all 10 signals.

Formally, let U_i denote a binary variable indicating whether an agent is underconfident ($U_i = 1$) or not ($U_i = 0$). The measured treatment effect of a good signal can be

decomposed into

$$\begin{aligned} \underbrace{\hat{\mathbb{E}}[Y_i|W_i = 1, X_i = 1, U_i] - \hat{\mathbb{E}}[Y_i|W_i = 0, X_i = 1, U_i]}_{\text{observed difference between treatment and control}} &= \underbrace{\mathbb{E}[Y_i(1) - Y_i(0)|W_i = 1, X_i = 1, U_i]}_{\text{treatment effect}} \\ &+ \underbrace{\mathbb{E}[Y_i(0)|W_i = 1, X_i = 1, U_i] - \mathbb{E}[Y_i(0)|W_i = 0, X_i = 1, U_i]}_{\text{selection}} \end{aligned}$$

The selection arises if participants in the treatment and control conditions behave differently (on average) even absent any intervention. Let's decompose the selection term further. Let $Pr(X_i = 1|W_i = 1, U_i = 1)$ denote the probability that a subject i observes a good signal, while receiving the treatment and being underconfident. We can expand the selection term as follows

$$\begin{aligned} \mathbb{E}[Y_i(0)|W_i = 1, X_i = 1, U_i] &= \mathbb{E}[Y_i(0)|W_i = 1, X_i = 1, U_i = 1] Pr(X_i = 1|W_i = 1, U_i = 1) \\ &+ \mathbb{E}[Y_i(0)|W_i = 1, X_i = 1, U_i = 0] (1 - Pr(X_i = 1|W_i = 1, U_i = 1)) \end{aligned}$$

and

$$\begin{aligned} \mathbb{E}[Y_i(0)|W_i = 0, X_i = 1, U_i] &= \mathbb{E}[Y_i(0)|W_i = 0, X_i = 1, U_i = 1] Pr(X_i = 1|W_i = 0, U_i = 1) \\ &+ \mathbb{E}[Y_i(0)|W_i = 0, X_i = 1, U_i = 0] (1 - Pr(X_i = 1|W_i = 0, U_i = 1)) \end{aligned}$$

Due to the random assignment to the treatment and control, it follows that

$$\gamma_1 \equiv \mathbb{E}[Y_i(0)|W_i = 1, X_i = 1, U_i = 1] = \mathbb{E}[Y_i(0)|W_i = 0, X_i = 1, U_i = 1]$$

and

$$\gamma_0 \equiv \mathbb{E}[Y_i(0)|W_i = 1, X_i = 1, U_i = 0] = \mathbb{E}[Y_i(0)|W_i = 0, X_i = 1, U_i = 0].$$

In other words, conditional on X, U , the assignment is as good as random. This simplifies the selection term

$$\begin{aligned} & \mathbb{E}[Y_i(0)|W_i = 1, X_i = 1, U_i] - \mathbb{E}[Y_i(0)|W_i = 0, X_i = 1, U_i] \\ &= \gamma_1 Pr(X_i = 1|W_i = 1, U_i = 1) + \gamma_0 (1 - Pr(X_i = 1|W_i = 1, U_i = 1)) \\ & \quad - \gamma_1 Pr(X_i = 1|W_i = 0, U_i = 1) - \gamma_0 (1 - Pr(X_i = 1|W_i = 0, U_i = 1)) \end{aligned}$$

This means that the selection term is zero if and only if participants in the treatment condition are as likely to receive a good signal as participants in the control condition,

$$Pr(U_i = 1|W_i = 1, X_i = 1) = Pr(U_i = 1|W_i = 0, X_i = 1).$$

However, participants in the control condition see all signals, while participants in the treatment see their own rank with probability $\frac{1}{2}$. Although we make sure that the participants are randomly allocated to the two groups, we cannot ensure that the signals relative to prior beliefs are randomly allocated to participants.

D.2 Solution: Matching Estimator

To deal with the potential selection issues (as well as to deal with potentially complex non-linearity of $\mu_w(\mathbf{x})$), I follow Heckman et al. (1998) and construct a matching estimator. For every participant in the treatment condition, I construct a counterfactual outcome based on decisions of participants in the control conditions with similar characteristics. In the example described in the previous section, we would match the underconfident agent who saw a good signal in the treatment condition with participants in the control condition, who were also underconfident and considered the same signal.

Following Heckman et al. (1998), I use a kernel regression to estimate the counterfactual outcomes $Y_i(0)$ for every participant in the treatment condition. The key identification assumption is that, conditional on all observables included in the matching procedure, the assignment of signals to participants is as good as random.

Formally, the treatment effect can be written as

$$\hat{\tau}(\mathbf{x}) = |N_T|^{-1} \sum_{i \in N_T} \left(Y_i - \sum_{j \in N_C} w_j^i (Y_j + \hat{\mu}_1(\mathbf{X}_i) - \hat{\mu}_1(\mathbf{X}_j)) \right), \quad (7)$$

where for each participant i in the treatment condition, w_j^i is a weight that I assign to a subject j from the control condition. The more similar participants i and j are (in terms of characteristics in \mathbf{X}), the higher the weight w_j^i . The weights are normalized such that $\sum_j w_j^i = 1$ and $w_j^i > 0, \forall_{i,j}$. The correction term $\hat{\mu}_1(\mathbf{X}_i) - \hat{\mu}_1(\mathbf{X}_j)$ removes the potential asymptotic bias of the matching estimator, as suggested by Abadie and Imbens (2011), where $\hat{\mu}_w(\mathbf{X})$ is a consistent regression estimator of $\mu_w(\mathbf{X})$. The intuition behind the bias correction is as follows. For a good match, the distance between \mathbf{X}_i and \mathbf{X}_j is small and the correction term $\hat{\mu}_1(\mathbf{X}_i) - \hat{\mu}_1(\mathbf{X}_j)$ vanishes. At the same time, the bias correction provides insurance in case of an imprecise match. In this case, the decision that a subject i would have made in the control condition becomes

$$\hat{Y}_i(0) = \hat{\mathbb{E}}[Y_i | W_i = 1, \mathbf{X}_i] + \sum_{j \in N_C} w_j^i \left(Y_j - \hat{\mathbb{E}}[Y_i | W_i = 1, \mathbf{X}_j] \right).$$

That is, the counterfactual outcome is equal to the regression prediction augmented with the matching term. The latter makes the whole estimator robust to a potential misspecification of the regression function stemming from the non-linearity of $\mu_w(x)$.

The weights w_j^i were constructed using the Epanechnikov kernel $K_{\mathbf{h}}(x)$ with a vector of parameters $\mathbf{h} > 0$

$$K_{\mathbf{h}}(\|\mathbf{X}_i - \mathbf{X}_j\|) = \begin{cases} \frac{3}{4} \left(1 - (\|\mathbf{h}(\mathbf{X}_i - \mathbf{X}_j)\|)^2 \right) & \text{if } \|\mathbf{h}(\mathbf{X}_i - \mathbf{X}_j)\| \leq 1 \\ 0 & \text{otherwise,} \end{cases} \quad (8)$$

where the length of \mathbf{h} is equal to the number of characteristics included in \mathbf{X} .

Thereby, the weights are given by

$$w_j^i(\mathbf{h}) = \begin{cases} \frac{K_{\mathbf{h}}(\|\mathbf{X}_i - \mathbf{X}_j\|)}{\sum_{j \in \mathcal{J}_i} K_{\mathbf{h}}(\|\mathbf{X}_i - \mathbf{X}_j\|)} & \text{if } j \in \mathcal{J}_i \\ 0 & \text{otherwise,} \end{cases} \quad (9)$$

where $\mathcal{J}_i \subset N_C$ is the set of individuals in the control condition who considered the same number as subject i . Epanechnikov kernel has similar interpretation to the nearest-neighbor matching often employed in the applied literature. It assigns positive weights only to a compact subset of neighbors whose characteristics are the closest to those of the target point. However, in contrast to nearest-neighbor matching which gives all the points in the neighborhood equal weight, Epanechnikov kernel assigns weights that decline smoothly with distance from the target point. This ensures that the resulting approximation of the conditional expectations $\mu_w(\mathbf{x})$ is smooth.

Parameter vector \mathbf{h} controls the support of the kernel – how many closest neighbors to include and how quickly the weights decay with the distance. I follow a standard practice in the literature and estimate \mathbf{h} using leave-one-out cross validation on the sample of participants in the control condition (see e.g. Hastie et al., 2009, Chapter 6). For a given \mathbf{h} , I estimate, using kernel regression, the probability each individual $k \in N_C$ assigned to Box 2

$$\hat{Y}_k(\mathbf{h}) = \sum_{j \in N_C \setminus \{k\}} w_j^k(\mathbf{h}) (Y_j + \hat{\mu}_1(\mathbf{X}_i) - \hat{\mu}_1(\mathbf{X}_j)).$$

I choose \mathbf{h} to minimize the mean squared prediction error

$$|N_C|^{-1} \sum_{k \in N_C} \left(Y_k - \hat{Y}_k(\mathbf{h}) \right)^2.$$

As for the choice of observables \mathbf{X} , I match agents in the treatment to those in the control who make decision regarding *the same* signal. Given this initial selection, I consider three specifications. In the baseline specification, I match participants based on their rank and prior belief distribution.

In the second specification, I match participants based on their rank and prior beliefs with respect to the number displayed on the computer screen. For example, if a participant observed the number “3” displayed on his screen, he would be matched based on his rank and how many points he allocated to rank 3 in the prior beliefs elicitation. In the third specification, I use only the prior belief distribution.

I estimate a consistent regression estimate $\hat{\mu}_w(x)$ using the estimator proposed in Hahn (1998)

$$\hat{\mu}_1(\mathbf{x}) = \frac{\hat{E}[YW|X = \mathbf{x}]}{\hat{E}[W|X = \mathbf{x}]} \quad (10)$$

To obtain consistent estimators for the conditional expectations on the right-hand side of (10), I use an OLS for the nominator and a logit regression for the denominator.

D.3 Inference

For a robust inference in a small sample, I employ the inferential techniques proposed in Abadie et al. (2010). The idea behind these techniques is to test whether the estimated treatment effect is large relative to the distribution of so-called “placebo effects”. The placebo effect is estimated by assigning the treatment status to a random sample of all participants and conducting the same regression analysis.

To this end, I draw a random sample of 160 observations (i.e. equal to the number of subjects in the treatment condition) from all observations in the experiment and assign them the treatment status. I follow the matching procedure to create a counterfactual for each of those 160 observations. Next, I use the observations and their counterfactuals to estimate a placebo treatment effect. I run the same regression as those described in Table 4 and store the resulting coefficients.

I repeat this procedure 20 000 times to obtain a distribution of the placebo effects. Those coefficients are presented in the histograms in Figures 10 and 11. The empirical distribution of the placebo effects allows me to calculate the p-values of a two-sided test to assess the statistical significance of the actual treatment effect.

E Additional Results

Table 12: The effect of the signal's valence if we control for how far the signal was from the subject's median belief.

	(1)	(2)	(3)	(4)	(5)
Good Signal	9.23** (4.16)	18.34*** (6.30)	18.42*** (6.14)	18.03*** (6.13)	25.19*** (6.67)
Distance	-0.88 (1.00)	0.60 (1.26)	3.78** (1.61)	2.53 (1.83)	2.53 (1.80)
Distance \times Good		-3.92* (2.05)	-3.88* (2.00)	-0.83 (2.91)	-8.73** (4.27)
Outside Priors			-17.37*** (5.70)	-10.54 (7.41)	-10.54 (7.29)
Outside Priors \times Good				-16.53 (11.54)	-38.35*** (14.31)
Good \times Distance \times Outside					11.85** (4.74)
Constant	3.26 (4.14)	-1.55 (4.82)	-2.02 (4.70)	-1.83 (4.68)	-1.83 (4.60)
Observations	160	160	160	160	160

Standard errors in parentheses

* $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$

Note: The dependent variable is the difference between numbers of points allocated to Box 2 in the Treatment and the Control (Kernel-Based Matching). "Distance" indicates the absolute difference between the signal and the subject's median belief.

F Matching

In this section, I present the results of my analysis described in Section 4.4 with the counterfactual based on different matching criteria. In Specification 2, I match participants based on their true rank and prior beliefs about the number under consideration (as opposed to the entire belief distribution in Specification 1). In Specification 3, I use only the prior belief distribution.

F.1 Matching Specification 2

Figure 15: Mean deviation from Bayes in the treatment condition and counterfactual.




Table 13: The effect of the signal's valence.

	(1)	(2)	(3)	(4)
Good Signal		11.44*** (4.13)	7.38* (4.12)	14.69*** (5.31)
Outside Priors			-15.18*** (4.10)	-8.13 (5.22)
Outside Priors \times Good				-17.82** (8.29)
Constant	7.49*** (2.09)	2.56 (2.71)	11.24*** (3.51)	7.21* (3.94)
Observations	160	160	160	160

Standard errors in parentheses

* $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$

Note: The dependent variable is the difference between numbers of points allocated to Box 2 in the treatment and in the counterfactual.

Figure 16: Distribution of coefficients at the Good Signal variable (Specification 4).



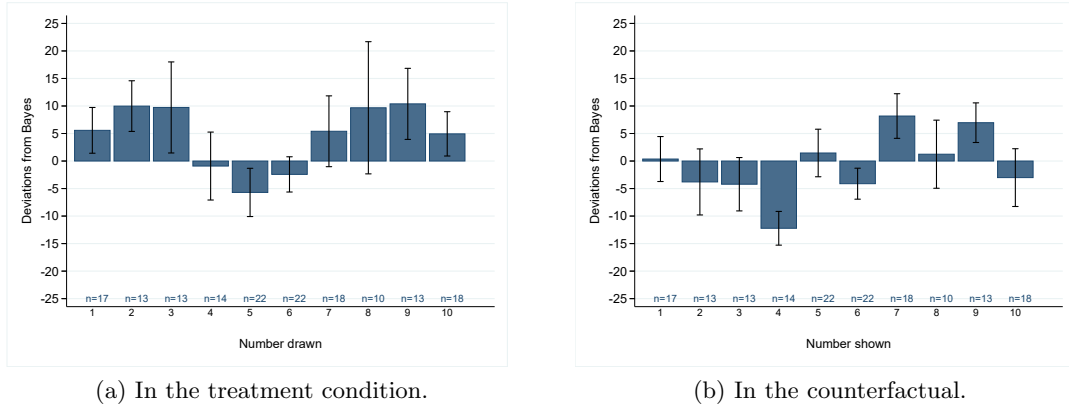


Figure 17: Distribution of coefficients at the interaction term (Specification 4).



F.2 Matching Specification 3

Figure 18: Mean deviation from Bayes in the treatment condition and counterfactual.



(a) In the treatment condition.

(b) In the counterfactual.

Table 14: The effect of the signal's valence.

	(1)	(2)	(3)	(4)
Good Signal		10.57*** (3.77)	8.88** (3.89)	15.88*** (5.01)
Outside Priors			-6.36 (3.87)	0.41 (4.92)
Outside Priors \times Good				-17.09** (7.82)
Constant	4.43** (1.91)	-0.13 (2.48)	3.51 (3.31)	-0.36 (3.72)
Observations	160	160	160	160

Standard errors in parentheses

* $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$

Note: The dependent variable is the difference between numbers of points allocated to Box 2 in the treatment and in the counterfactual (kernel-based matching). "Good Signal" indicator variable takes value 1 if the signal was below or equal to the median of subject's belief distribution, and 0 otherwise. "Outside priors" indicator variable takes value 1 if the subject attached a zero prior probability to the signal being his rank, and 0 otherwise.

Figure 19: Distribution of coefficients at the Good Signal variable (Specification 4).





Figure 20: Distribution of coefficients at the interaction term (Specification 4).



G Overconfidence

Table 15: Belief distribution and Confidence Type.

	All	Overconfident	Unbiased	Underconfident
Rank	5.65 (2.69)	7.13 (1.97)	4.38 (1.44)	2.93 (1.98)
Mean	4.47 (1.75)	4.17 (1.53)	4.43 (1.35)	5.13 (2.16)
1st Quartile	3.71 (1.74)	3.40 (1.47)	3.71 (1.26)	4.41 (2.21)
Median	4.45 (1.79)	4.14 (1.54)	4.38 (1.44)	5.16 (2.20)
3rd Quartile	5.16 (1.87)	4.84 (1.70)	5.21 (1.44)	5.84 (2.19)
Range	4.89 (1.57)	4.82 (1.64)	5.04 (1.44)	5.00 (1.60)
Bias	2.44 (1.84)	2.99 (1.74)	0 (.)	2.23 (1.58)
N	209	127	24	58

G.1 The Benoit-Dubra Critique

(soon to be updated)

H RTC Registration details

(soon to be updated)

I Information Acquisition

(soon to be updated)

J Instructions

(soon to be updated)